

000 001 002 003 004 005 A WALKTHROUGH OF NANOCHEAT 006 007 008 009 010 011 012 013 014 015 016 017 018 019 020 021 022 023 024 025 026 027 028 029 030 031 032 033 034 035 036 037 038 039 040 041 042 043 044 045 046 047 048 049 050 051 052 053

Anonymous authors
Paper under double-blind review

ABSTRACT

This walkthrough of nanochat aims to study a complete, modern LLM implementation with cutting-edge techniques, all in a clean, accessible codebase. It is also the perfect fit to learn about the complete pipeline from tokenization and pretraining to finetuning and RL.

1 QUICKSTART

The Quickstart of nanochat demonstrates a depth-20 speedrun model.

It is a GPT-like Llama model with a context size of up to 65,536 characters, 1280 hidden dimensions, and it is a 20-layer Transformer with 10 heads in each layer.

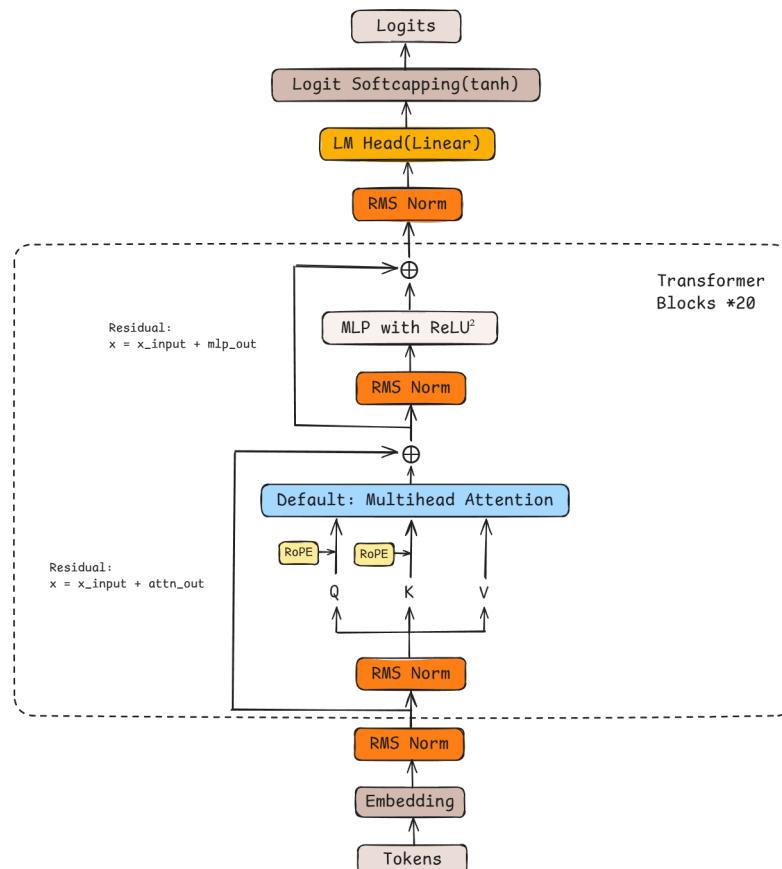


Figure 1: Architecture

054 1.1 SIZE
055056 For this depth-20 model:
057

- 058 • $L = \text{Layers} = 20$
- 059 • $D = \text{Model Dimension(Hidden Dimension)} = 20 \times 64 = 1280$
- 060 • $V = \text{Vocabulary size} = 65,536$

061
062 **Universal Formula:**
063

$$\begin{aligned}\text{Parameters} &= V \times D + L \times (4 \times D^2 + 2 \times (D \times (4 \times D))) + V \times D \\ &= V \times D + 20 \times (4 \times D^2 + 2 \times (D \times 4D)) + V \times D \\ &= 83,886,080 + 20 \times (6,553,600 + 13,107,200) + 83,886,080 \\ &= 560,988,160\end{aligned}$$

- 069 • **For embedding and unembedding:** each has $V \times D$ parameters
070
- 071 • **For each Transformer block:**
 - 072 – **Attention Layer:** 4 weight matrices(Q, K, V and ouput projection), each has D^2
073 parameters
 - 074 – **MLP Layer:** 2 weight matrices(4x expansion and projection), each has $D \times (4 \times D)$
075 parameters

076 1.2 ATTENTION HEADS
077

The d20 model has $\text{ceil}(D/128) = 10$ Attention Heads. The model applies Multi-Head Attention(MHA) as default for simplicity, but supports Grouped Query Attention(GQA)/Multi-Query Attention(MQA) as well.

Within the attention heads, the model introduces QK normalization, which stabilizes logit scale across layers/steps and reduces sensitivity to learning rate, giving rise to less softmax saturation, fewer exploding heads as well as smoother gradients.

085 1.3 POSITIONAL EMBEDDING
086

The model applies RoPE to Q, K within the attention heads instead of learned position embedding. The parameter-free RoPE not only improves efficiency, but also enhances stability in longer context, especially when it is combined with Q, K norm.(Su et al., 2023)

091 1.4 MLP WITH RELU²
092

The model employs the original MLP with squared ReLU activation function.

Compared with GLU variants which trade $1.5 \times$ parameters for 1-3% performance improvement(Shazeer, 2020), standard MLP saves up to 20% training cost and time. This configuration is particularly suitable for the 561M nanochat model, which aims to be "the best ChatGPT that \$100 can buy."

Compared with ReLU, SwiGLU, and ReGLU, ReLU² excels across all three evaluation metrics: the trade-off between sparsity and performance, the predictivity of sparsity, and the hardware affinity (Zhang et al., 2024).

102 1.5 RMSNORM
103

The model applies Pre-norm RMSNorm with no learnable parameters after token embedding, before attention and MLP in each block, and once more before the LM head.

Compared to LayerNorm, which does 5 operations each time, parameter-free RMSNorm does 2. Likewise, it saves training cost and time. Besides, modded-nanogpt(Jordan et al., 2024) experiments

108 shows that parameter-free RMSNorm trains just as well at small scale. We may hypothesize that at
 109 100M-1B parameter scale, the model has enough capacity that normalization does not need learnable
 110 adaptation.
 111

112 1.6 FLASHNORM 113

114 The combination of bias-free linear layers, QK normalization and parameter-free are in line with
 115 the concept of FlashNorm(Graef et al., 2025). The paper also introduces ways of optimization by
 116 integrating RoPE into normalization and computing RMSNorm and linear layers in parallel, but
 117 Karpathy did not adopt them.

118 1.7 UNTIED EMBEDDING/UNEMBEDDING 119

120 The untied embedding/unembedding let input embeddings and output classifier specialize independently,
 121 allowing for different learning rate configuration(0.2 for embedding and 0.0004 for unembedding)
 122 and different optimizers(AdamW for embeddings, Muon for transformer)(Karpathy, 2025).
 123

124 1.8 LOGITS SOFTCAP 125

126 The model employs Logits softcap(\tanh) to clip extreme logits, avoiding unstable gradients and
 127 overflow risk. According to (Rybakov et al., 2024), its combination with QK normalization improves
 128 training robustness and sampling temperature behavior.

129 130 ACKNOWLEDGMENTS

131 Thanks to Zhengfu He for guiding my research career :P.
 132

133 134 REFERENCES

- 135 Nils Graef, Andrew Wasilewski, and Matthew Clapp. Flashnorm: fast normalization for llms, 2025.
 136 URL <https://arxiv.org/abs/2407.09577>.
 137
- 138 Keller Jordan, Jeremy Bernstein, Brendan Rappazzo, @fernbear.bsky.social, Boza Vlado,
 139 You Jiacheng, Franz Cesista, Braden Koszarsky, and @Grad62304977. modded-nanogpt:
 140 Speedrunning the nanogpt baseline, 2024. URL <https://github.com/KellerJordan/modded-nanogpt>.
 141
- 142 Andrej Karpathy. nanochat: The best chatgpt that 100canbuy, 2025. URL.
 143
- 144 Oleg Rybakov, Mike Chrzanowski, Peter Dykas, Jinze Xue, and Ben Lanir. Methods of improving
 145 llm training stability, 2024. URL <https://arxiv.org/abs/2410.16682>.
- 146 Noam Shazeer. Glu variants improve transformer, 2020. URL <https://arxiv.org/abs/2002.05202>.
 147
- 149 Jianlin Su, Yu Lu, Shengfeng Pan, Ahmed Murtadha, Bo Wen, and Yunfeng Liu. Roformer: En-
 150 hanced transformer with rotary position embedding, 2023. URL <https://arxiv.org/abs/2104.09864>.
 151
- 152 Zhengyan Zhang, Yixin Song, Guanghui Yu, Xu Han, Yankai Lin, Chaojun Xiao, Chenyang Song,
 153 Zhiyuan Liu, Zeyu Mi, and Maosong Sun. Relu2 wins: Discovering efficient activation functions
 154 for sparse llms. *arXiv preprint arXiv:2402.03804*, 2024.

155
 156
 157
 158
 159
 160
 161